

COSC 520 Project

Breaking the Sorting Barrier

Hamza Muhammad Anwar

Bakdauren Narbayev

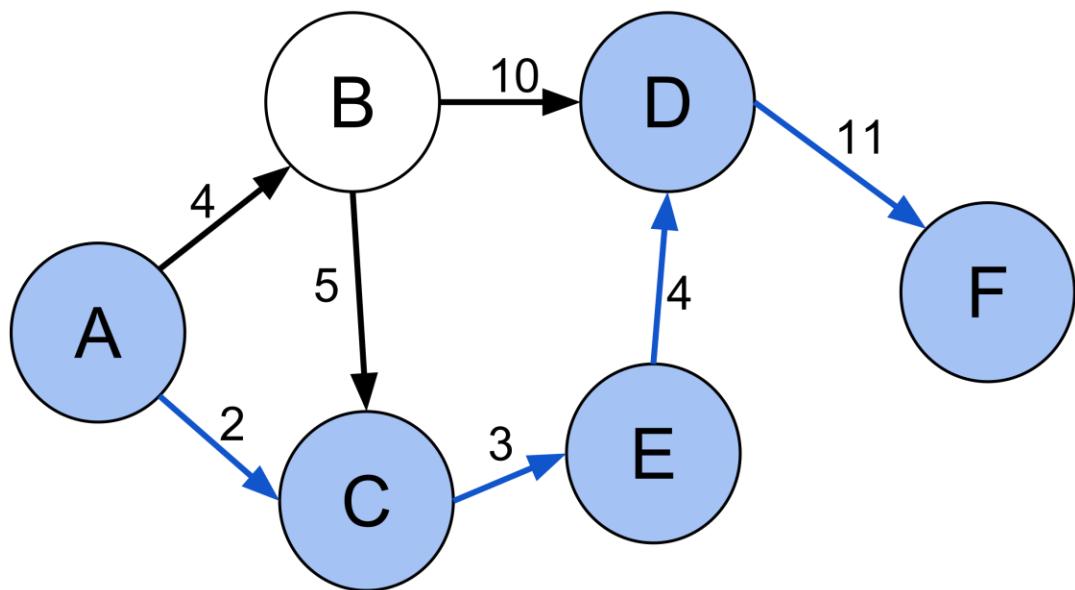
Ghaith Chrit

Overview

- Chose and implemented a new algorithm: BMSSP
- Benchmarked it with other similar algorithms
- Made a GUI to visualize their performance

Shortest Path Problem

- Finding the path between two vertices with minimum weights
- Directed, undirected, mixed



Motivation

- BMSSP: Do theoretical guarantees materialize during empirical comparisons?
- Beat a 66-year-old record
- Award-winning algorithm

Algorithms

- Bellman-Ford (Bellman, 1958)
- Dijkstra (Dijkstra, 1959)
- BMSSP (Duan et al. 2025)

Bellman-Ford Algorithm

- Slower than Dijkstra
- Can work with negative-edge weights
- Max $|V|-1$ edges in the path (No cycles)
- Time complexity: $O(VE)$
- Space complexity: $O(V)$

Bellman-Ford Algorithm Pseudocode

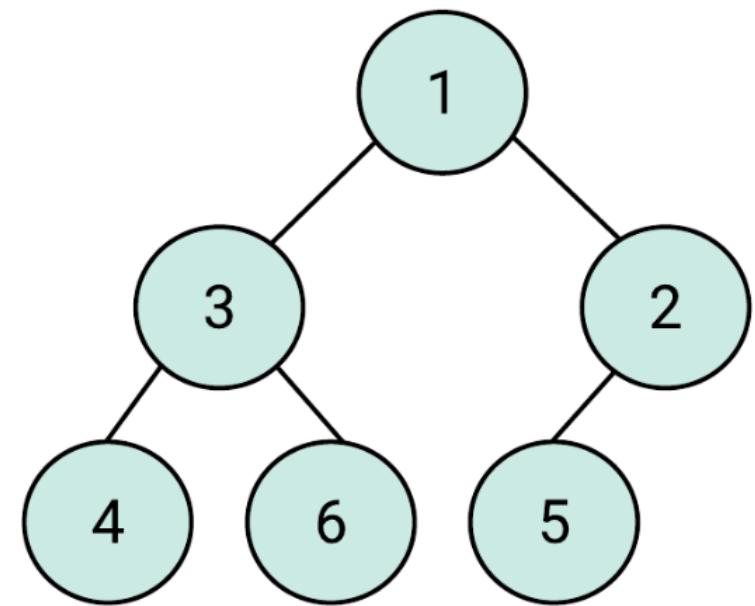
```
Initialize ()  
for  $i = 1$  to  $|V| - 1$   
    for each edge  $(u, v) \in E$ :  
        Relax( $u, v$ )  
    for each edge  $(u, v) \in E$   
        do if  $d[v] > d[u] + w(u, v)$   
            then report a negative-weight cycle exists
```

Dijkstra's Algorithm

- Cannot handle negative-edge weights
- Greedy
- Time complexity: $O((V+E) \log V)$
- Space complexity: $O(V)$

Dijkstra's Algorithm – Min Heap

- Priority queue
- Root node contains the smallest element
- All nodes contain elements less than or equal to their child nodes



Min heap

Dijkstra's Algorithm Pseudocode

```
1: procedure DIJKSTRA( $G, s$ )
2:    $V \leftarrow$  vertices of  $G$ 
3:    $E \leftarrow$  edges of  $G$ 
4:   for each vertex  $v \in V$  do                                ▷ Initialize distances and priority queue
5:      $d[v] \leftarrow \infty$ 
6:   end for
7:    $d[s] \leftarrow 0$ 
8:    $Q \leftarrow$  priority queue containing all vertices in  $V$ 
9:   while  $Q$  is not empty do                                ▷ Process vertices in order of distance
10:     $u \leftarrow$  vertex in  $Q$  with minimum  $d[u]$ 
11:    Remove  $u$  from  $Q$ 
12:    for each neighbor  $v$  of  $u$  do
13:       $alt \leftarrow d[u] + w(u, v)$ 
14:      if  $alt < d[v]$  then
15:         $d[v] \leftarrow alt$ 
16:        Update priority of  $v$  in  $Q$ 
17:      end if
18:    end for
19:  end while
20:  return  $d$ 
21: end procedure
```

Applications

- Network Routing
- Transportation and GIS
- Pathfinding for Autonomous Robots

Dijkstra's algorithm is still the fastest if the distances need to be obtained in increasing order!

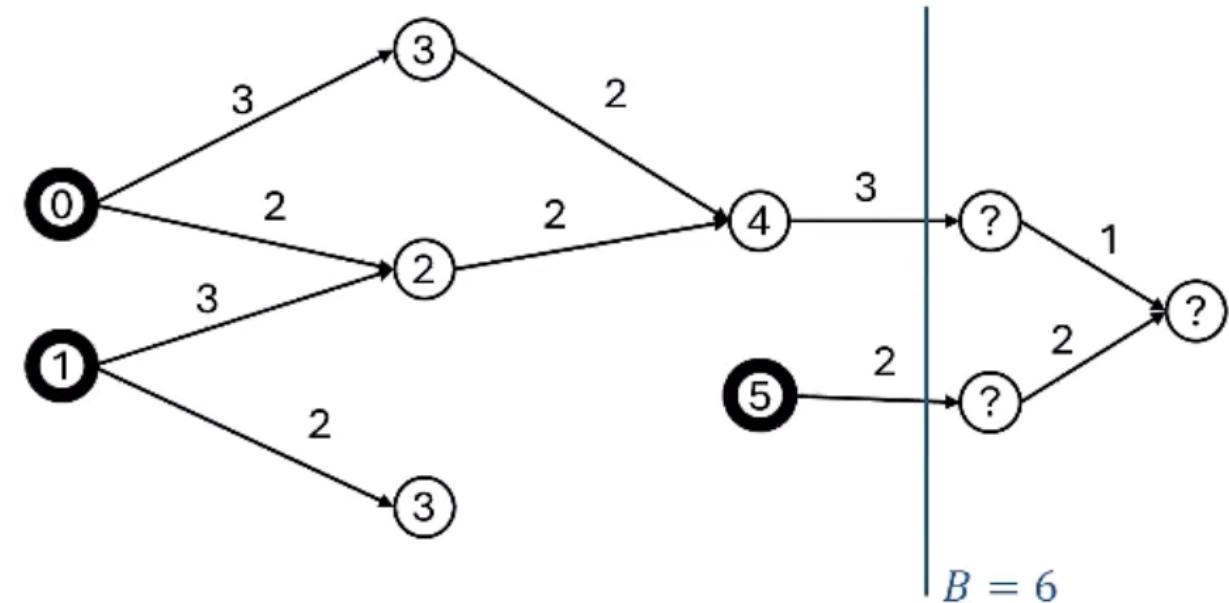
- Sparse directed graph with non-negative weights
- Time complexity: $O(E \log^{2/3} V)$ Beats Dijkstra's algorithm!
- Space complexity: $\Omega(V \log V^{1/3})$

Bounded Multi-Source Shortest Path (BMSSP)

BMSSP Problem

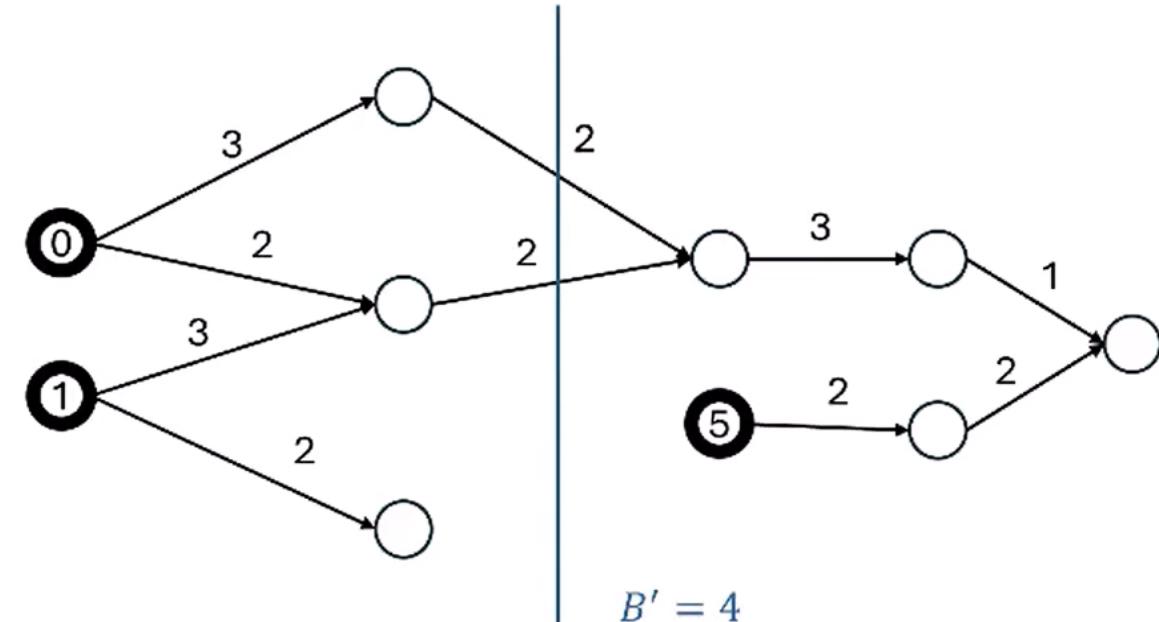
Given a set S of sources with initial distances, and an upper bound B , for every vertex x , suppose the shortest path from any vertex in S to x is $d[x]$.

The algorithm reports whether $d[x] < B$, and if $d[x] < B$, the algorithm finds $d[x]$.¹



Divide and Conquer

- Suppose we're given $B' \leq B$ such that exactly half of the vertices satisfy $d[x] < B'$.
- We can divide the problem into 2 equal halves:
 1. vertices with $d[x] < B'$
 2. vertices with $B' \leq d[x] < B$

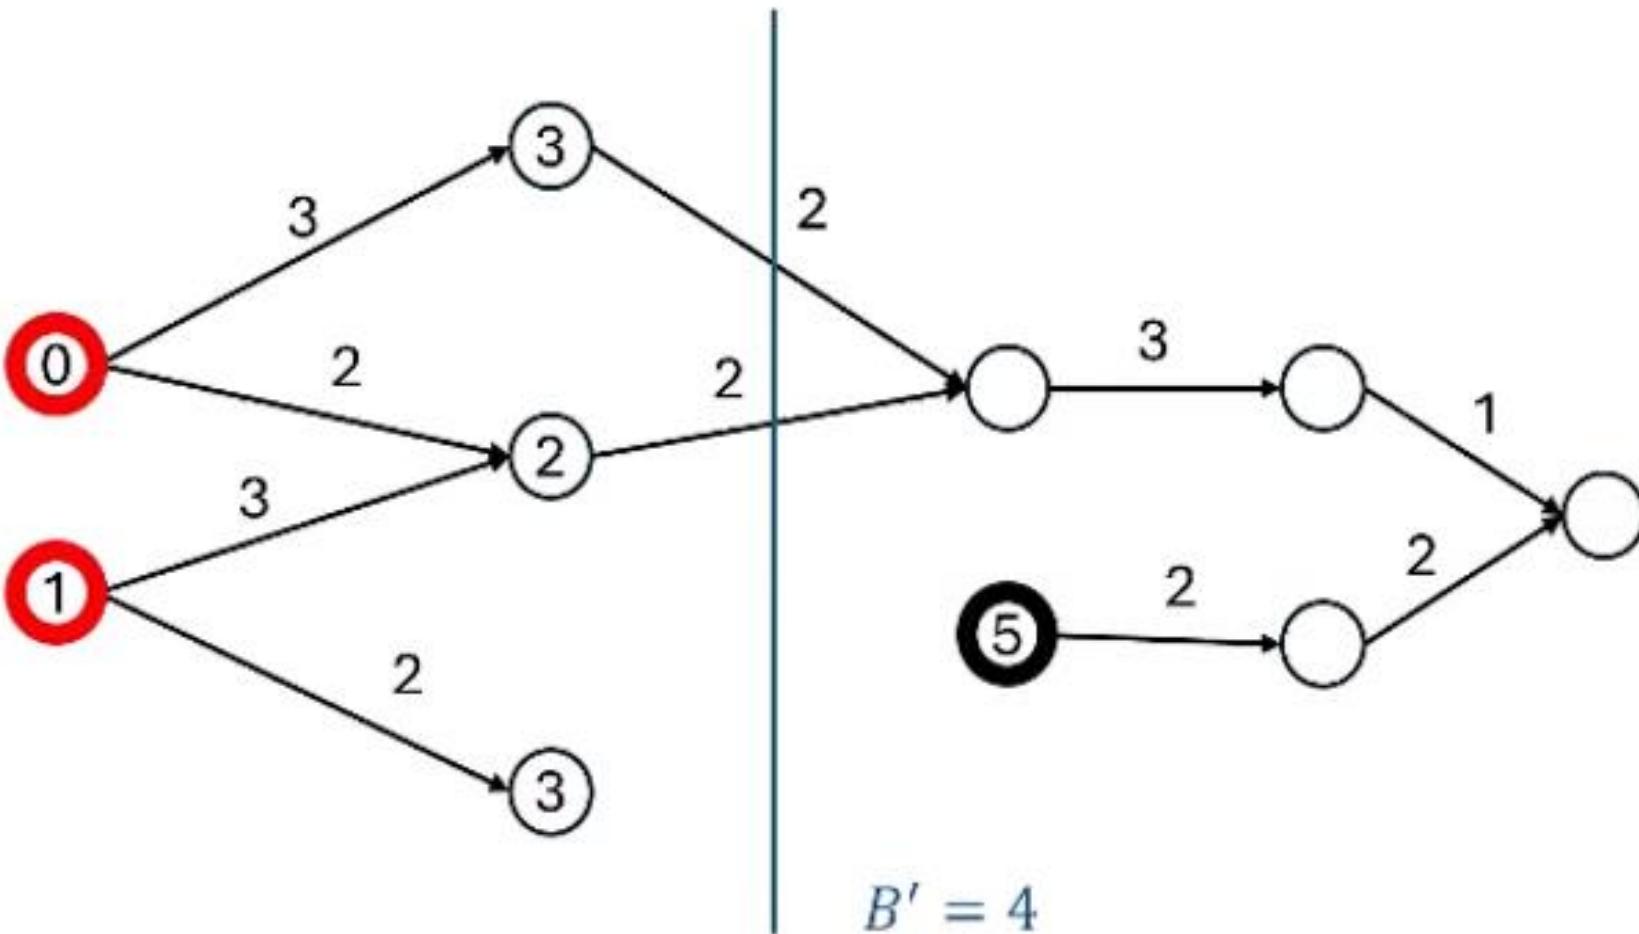


BMSSP Problem

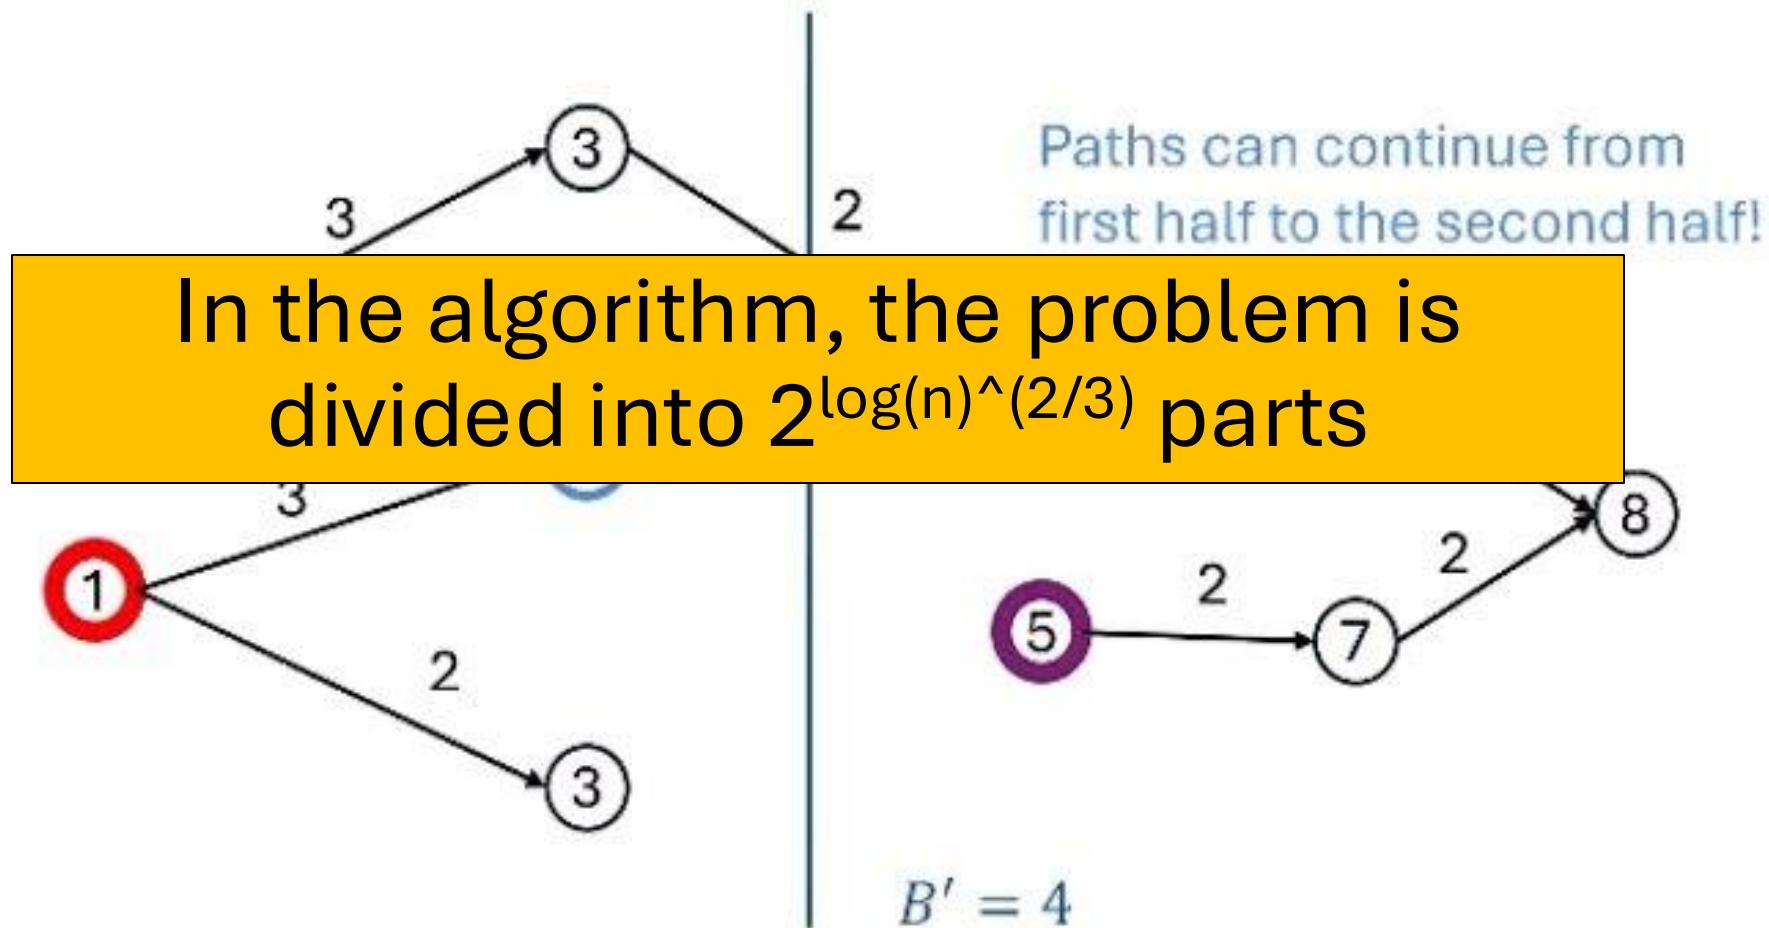
Given a set S of sources with initial distances, and an upper bound B , for every vertex x , suppose the shortest path from any vertex in S to x is $d[x]$.

The algorithm reports whether $d[x] < B$, and if $d[x] < B$, the algorithm finds $d[x]$.

Divide and Conquer (1st half)

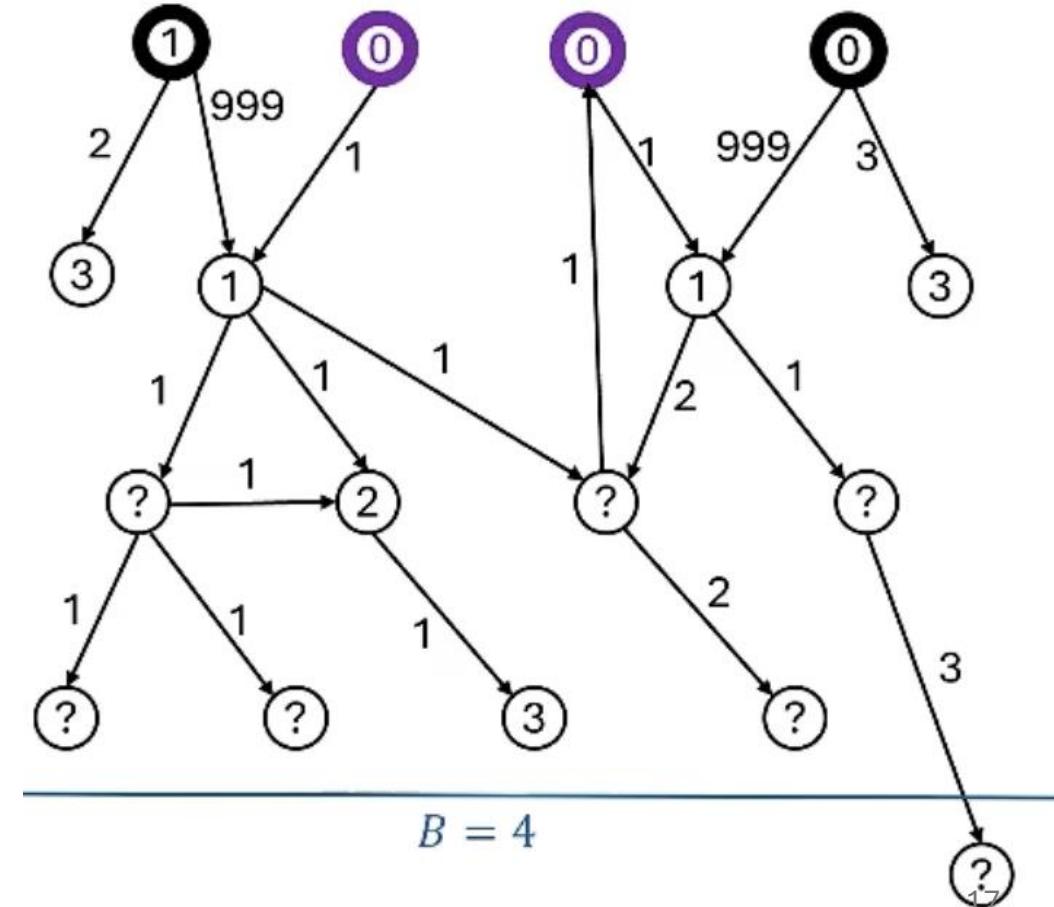


Divide and Conquer (2nd half)



Pivot Pruning

- The set of source vertices might contain all the vertices... **Too bad!**
- Need to choose vertices for a sub-problem carefully
- Let's "shrink" S by a factor of $\log^{1/3}(n)$



Algorithm 1 Finding Pivots

Piv 1: **function** FINDPIVOTS(B, S)

- R • **requirement:** for every incomplete vertex v with $d(v) < B$, the shortest path to v visits some complete vertex in S

- C • **returns:** sets P, W satisfying the conditions in Lemma 3.2

C 2: $W \leftarrow S$

tl 3: $W_0 \leftarrow S$

• T 4: **for** $i \leftarrow 1$ to k **do**

▷ Relax for k steps

O 5: $W_i \leftarrow \emptyset$

V 6: **for** all edges (u, v) with $u \in W_{i-1}$ **do**

7: **if** $\hat{d}[u] + w_{uv} \leq \hat{d}[v]$ **then**

8: $\hat{d}[v] \leftarrow \hat{d}[u] + w_{uv}$

9: **if** $\hat{d}[u] + w_{uv} < B$ **then**

10: $W_i \leftarrow W_i \cup \{v\}$

11: $W \leftarrow W \cup W_i$

12: **if** $|W| > k|S|$ **then**

13: $P \leftarrow S$

14: **return** P, W

15: $F \leftarrow \{(u, v) \in E : u, v \in W, \hat{d}[v] = \hat{d}[u] + w_{uv}\}$ ▷ F is a directed forest under Assumption 2.1

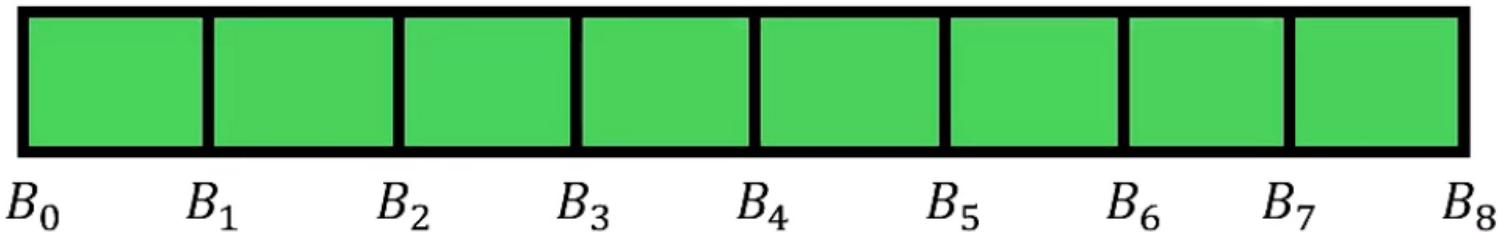
16: $P \leftarrow \{u \in S : u \text{ is a root of a tree with } \geq k \text{ vertices in } F\}$

17: **return** P, W

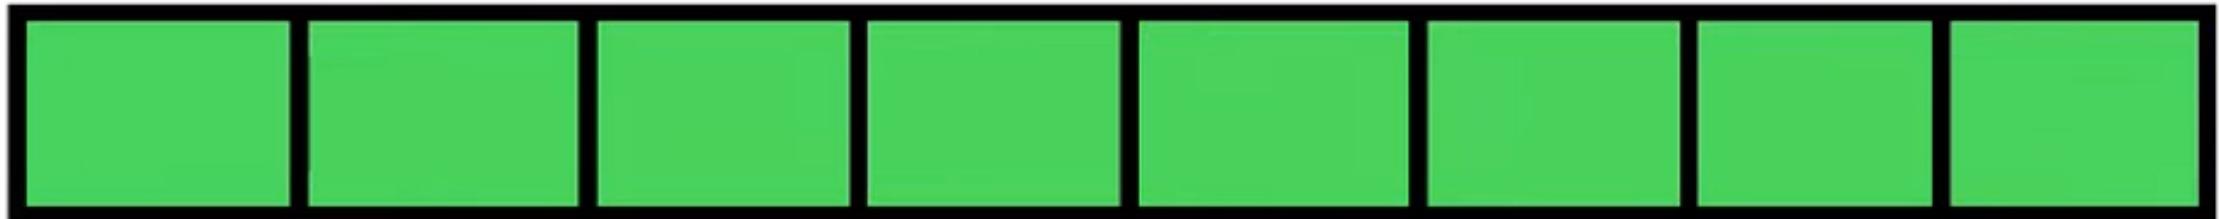
Short
in S



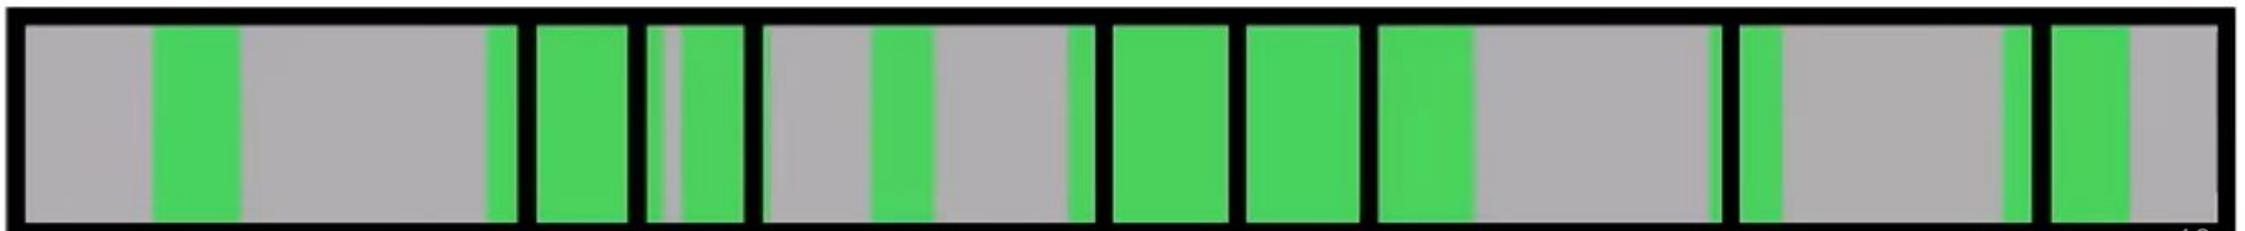
How to divide into sub-problems?



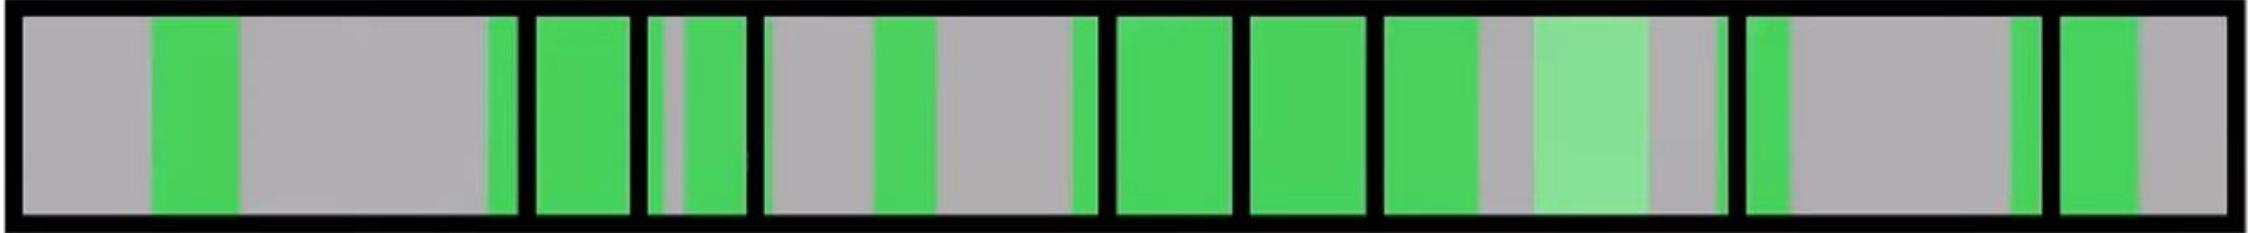
Ideally



Reality (guessing!)



What if a sub-problem becomes too large?

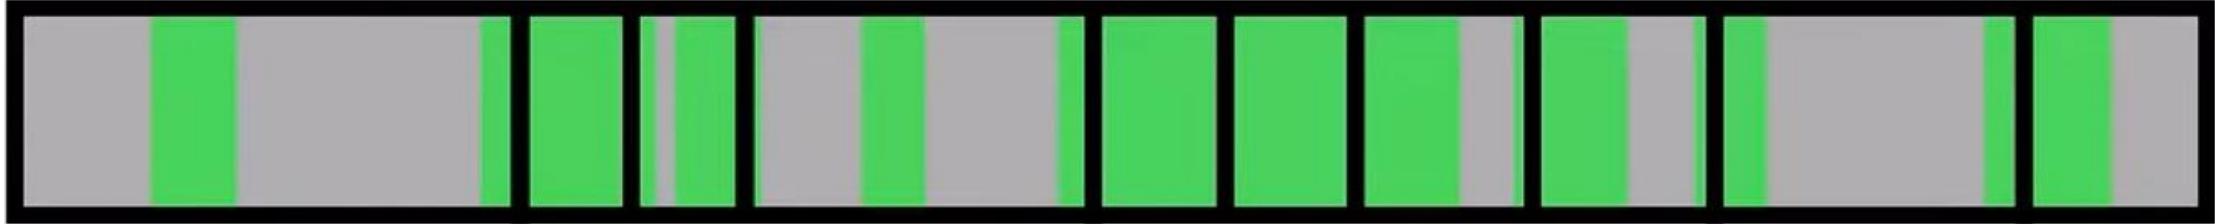


New distances
were computed

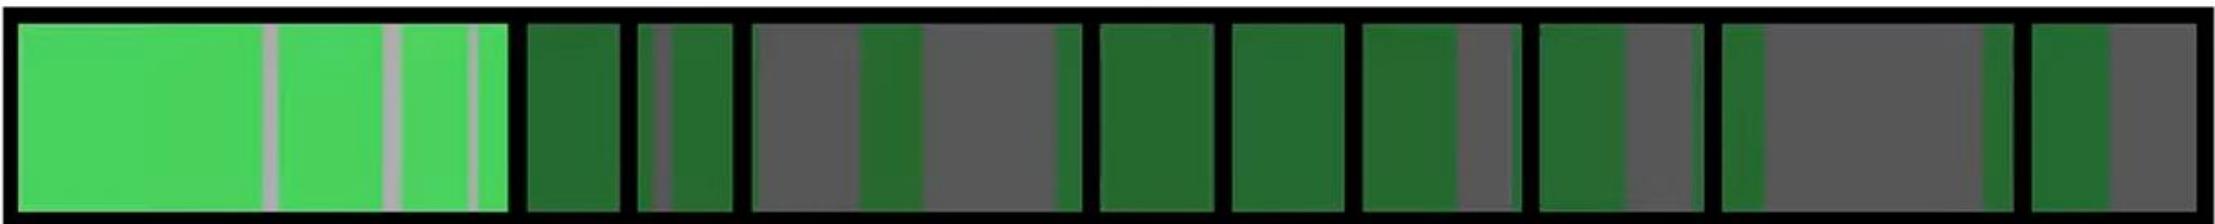


Split the
sub-problem

What if already started solving such sub-problem?

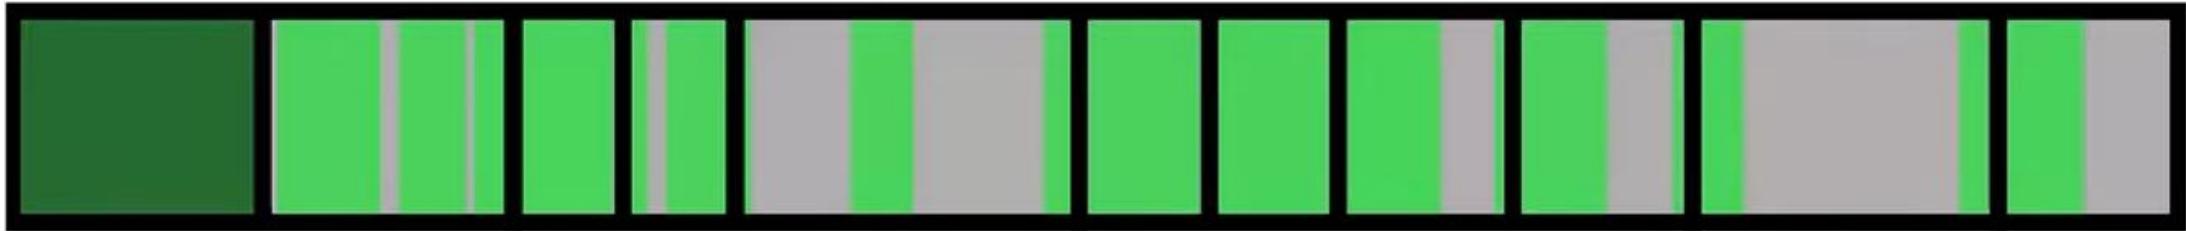


Started this sub-
problem

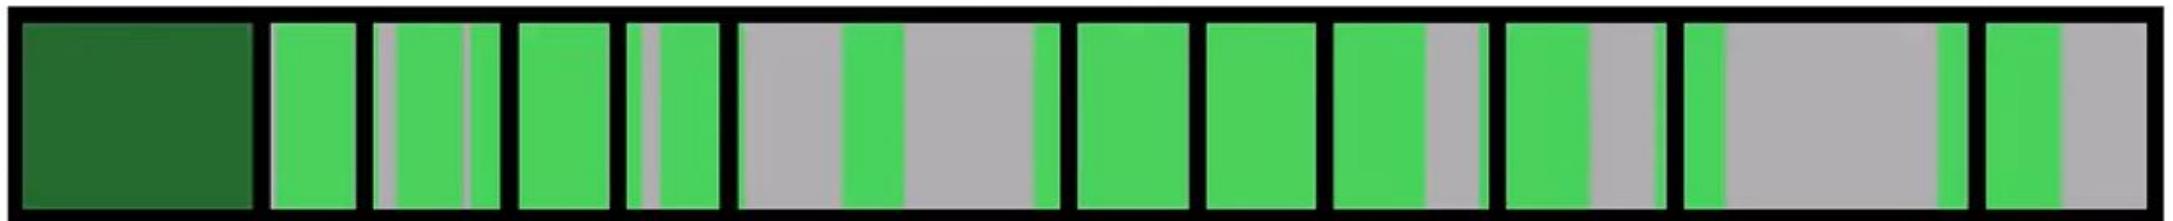


More distances
were computed

What if already started solving such sub-problem?

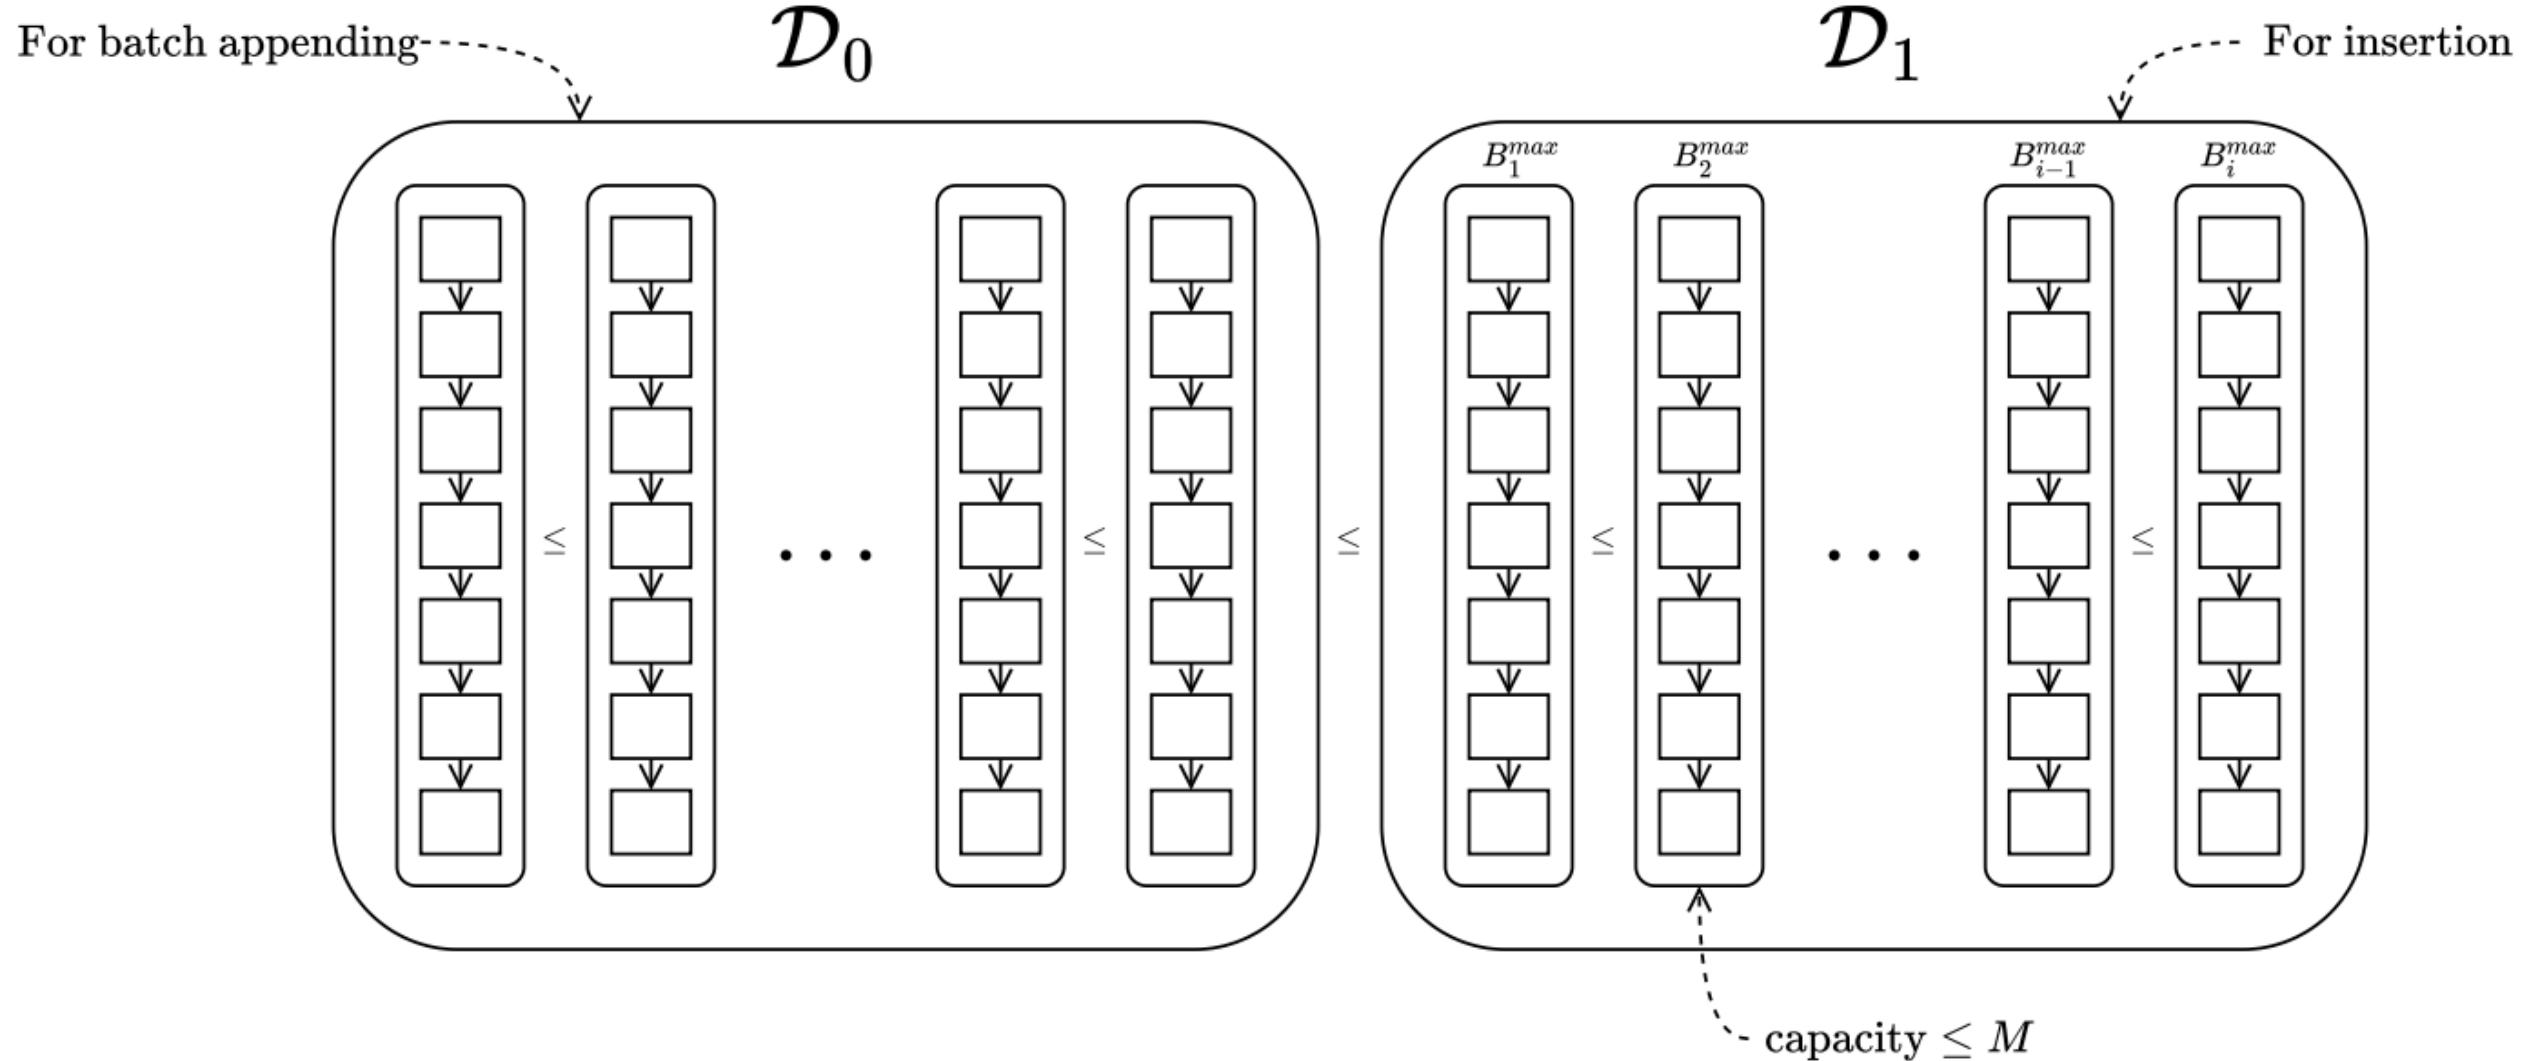


Go back to the
parent problem



Rebalance the
unfinished work

How are the bounds managed?



Base Case

Algorithm 2 Base Case of BMSSP

```
1: function BASECASE( $B, S$ )
   • requirement 1:  $S = \{x\}$  is a singleton, and  $x$  is complete
   • requirement 2: for every incomplete vertex  $v$  with  $d(v) < B$ , the shortest path to  $v$  visits  $x$ 
   • returns 1: a boundary  $B' \leq B$ 
   • returns 2: a set  $U$ 
2:    $U_0 \leftarrow S$ 
3:   initialize a binary heap  $\mathcal{H}$  with a single element  $\langle x, \hat{d}[x] \rangle$  [Wil64]
4:   while  $\mathcal{H}$  is non-empty and  $|U_0| < k + 1$  do
5:      $\langle u, \hat{d}[u] \rangle \leftarrow \mathcal{H}.\text{EXTRACTMIN}()$ 
6:      $U_0 \leftarrow U_0 \cup \{u\}$ 
7:     for edge  $e = (u, v)$  do
8:       if  $\hat{d}[u] + w_{uv} \leq \hat{d}[v]$  and  $\hat{d}[u] + w_{uv} < B$  then
9:          $\hat{d}[v] \leftarrow \hat{d}[u] + w_{uv}$ 
10:        if  $v$  is not in  $\mathcal{H}$  then
11:           $\mathcal{H}.\text{INSERT}(\langle v, \hat{d}[v] \rangle)$ 
12:        else
13:           $\mathcal{H}.\text{DECREASEKEY}(\langle v, \hat{d}[v] \rangle)$ 
14:        if  $|U_0| \leq k$  then
15:          return  $B' \leftarrow B, U \leftarrow U_0$ 
16:        else
17:          return  $B' \leftarrow \max_{v \in U_0} \hat{d}[v], U \leftarrow \{v \in U_0 : \hat{d}[v] < B'\}$ 
```

Running Dijkstra's
algorithm on singleton
source vertex

Recursive Case

Algorithm 3 Bounded Multi-Source Shortest Path

```

1: function BMSSP( $l, B, S$ )
   • requirement 1:  $|S| \leq 2^{ll}$ 
   • requirement 2: for every incomplete vertex  $x$  with  $d(x) < B$ , the shortest path to  $x$  visits some
     complete vertex  $y \in S$ 
   • returns 1: a boundary  $B' \leq B$ 
   • returns 2: a set  $U$ 
2:   if  $l = 0$  then
3:     return  $B', U \leftarrow \text{BASECASE}(B, S)$ 
4:    $P, W \leftarrow \text{FINDPIVOTS}(B, S)$ 
5:    $\mathcal{D}.\text{INITIALIZE}(M, B)$  with  $M = 2^{(l-1)l}$                                  $\triangleright \mathcal{D}$  is an instance of Lemma 3.3
6:    $\mathcal{D}.\text{INSERT}(\langle x, \hat{d}[x] \rangle)$  for  $x \in P$ 
7:    $i \leftarrow 0$ ;  $B'_0 \leftarrow \min_{x \in P} \hat{d}[x]$ ;  $U \leftarrow \emptyset$                                  $\triangleright$  If  $P = \emptyset$  set  $B'_0 \leftarrow B$ 
8:   while  $|U| < k2^{ll}$  and  $\mathcal{D}$  is non-empty do
9:      $i \leftarrow i + 1$ 
10:     $B_i, S_i \leftarrow \mathcal{D}.\text{PULL}()$ 
11:     $B'_i, U_i \leftarrow \text{BMSSP}(l - 1, B_i, S_i)$ 
12:     $U \leftarrow U \cup U_i$ 
13:     $K \leftarrow \emptyset$ 
14:    for edge  $e = (u, v)$  where  $u \in U_i$  do
15:      if  $\hat{d}[u] + w_{uv} \leq \hat{d}[v]$  then
16:         $\hat{d}[v] \leftarrow \hat{d}[u] + w_{uv}$ 
17:        if  $\hat{d}[u] + w_{uv} \in [B_i, B)$  then
18:           $\mathcal{D}.\text{INSERT}(\langle v, \hat{d}[u] + w_{uv} \rangle)$ 
19:        else if  $\hat{d}[u] + w_{uv} \in [B'_i, B_i)$  then
20:           $K \leftarrow K \cup \{\langle v, \hat{d}[u] + w_{uv} \rangle\}$ 
21:         $\mathcal{D}.\text{BATCHPREPEND}(K \cup \{\langle x, \hat{d}[x] \rangle : x \in S_i \text{ and } \hat{d}[x] \in [B'_i, B_i)\})$ 
22:    return  $B' \leftarrow \min\{B'_i, B\}$ ;  $U \leftarrow U \cup \{x \in W : \hat{d}[x] < B'\}$ 

```

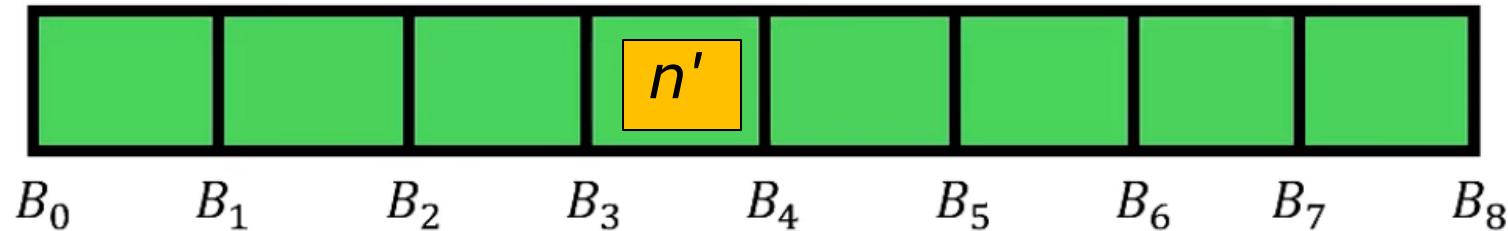
Intuition for Time Complexity

Per vertex:

- Operations in the data structure take $O(1)$ time
- Might be in pivot pruning: $O(k^2)$
- Might be in the base case: $O(1)$ (constant degree assumption)

Total: $O(nk^2) = O(n \log^{2/3}(n))$

Intuition for Space Complexity



Per sub-problem: $\Omega(n')$

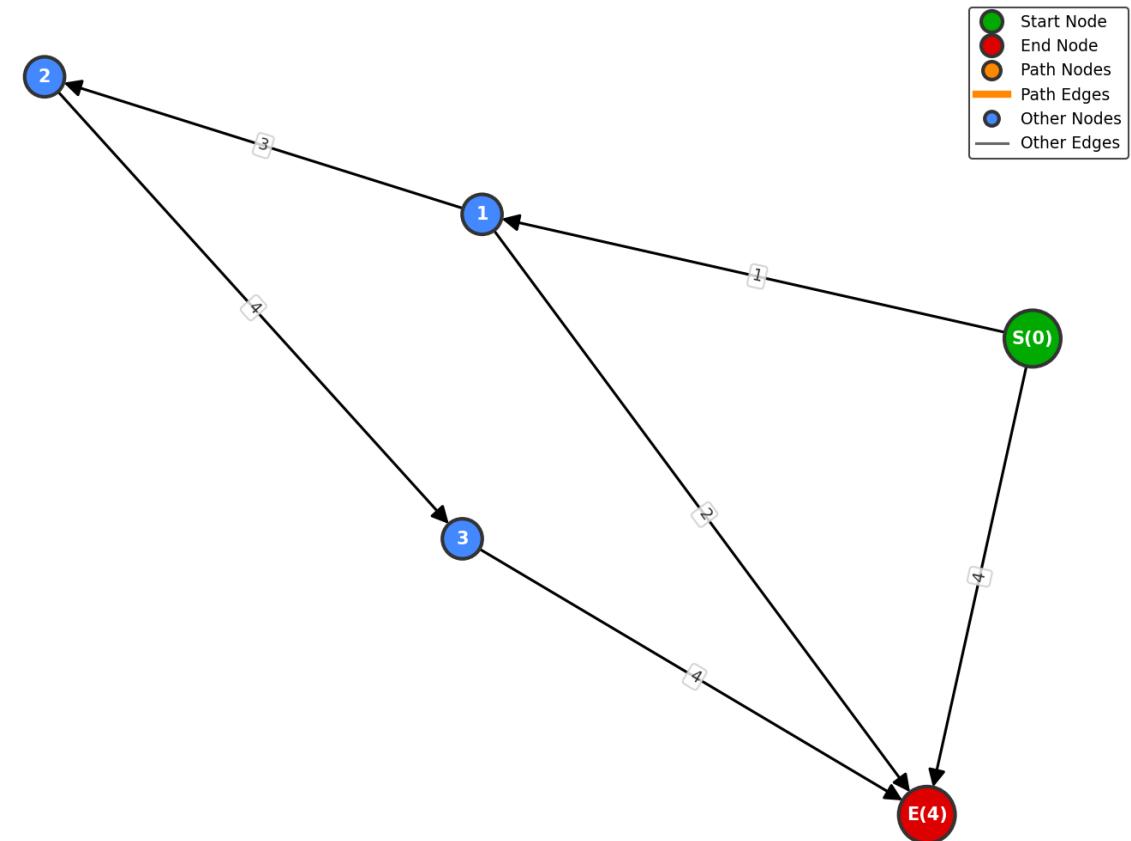
Per level: $\Omega(n)$

Total: $\Omega(n \log(n) / t) = \Omega(n \log^{1/3}(n))$

Much worse in our implementation :(

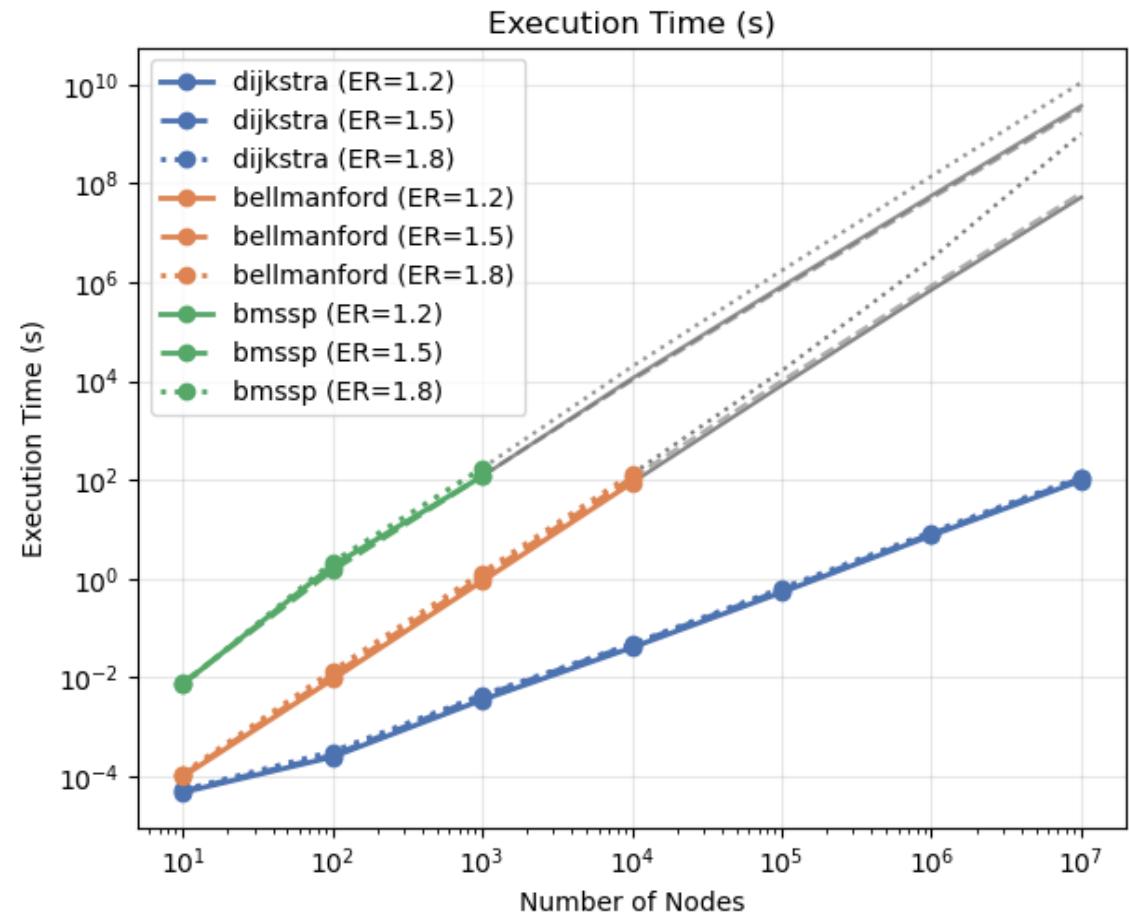
Dataset Generation

- Limited our study to directed sparse-graphs.
- Created a synthetic dataset with specific number of nodes and varying edge ratios.
- Experiments ran for:
 - 10 to 10^7 nodes
 - 1.2, 1.5, 1.8 edge ratios
- Weights were uniformly sampled from 1 to 10.
- Controlled in- and out-degree of nodes during benchmarking.



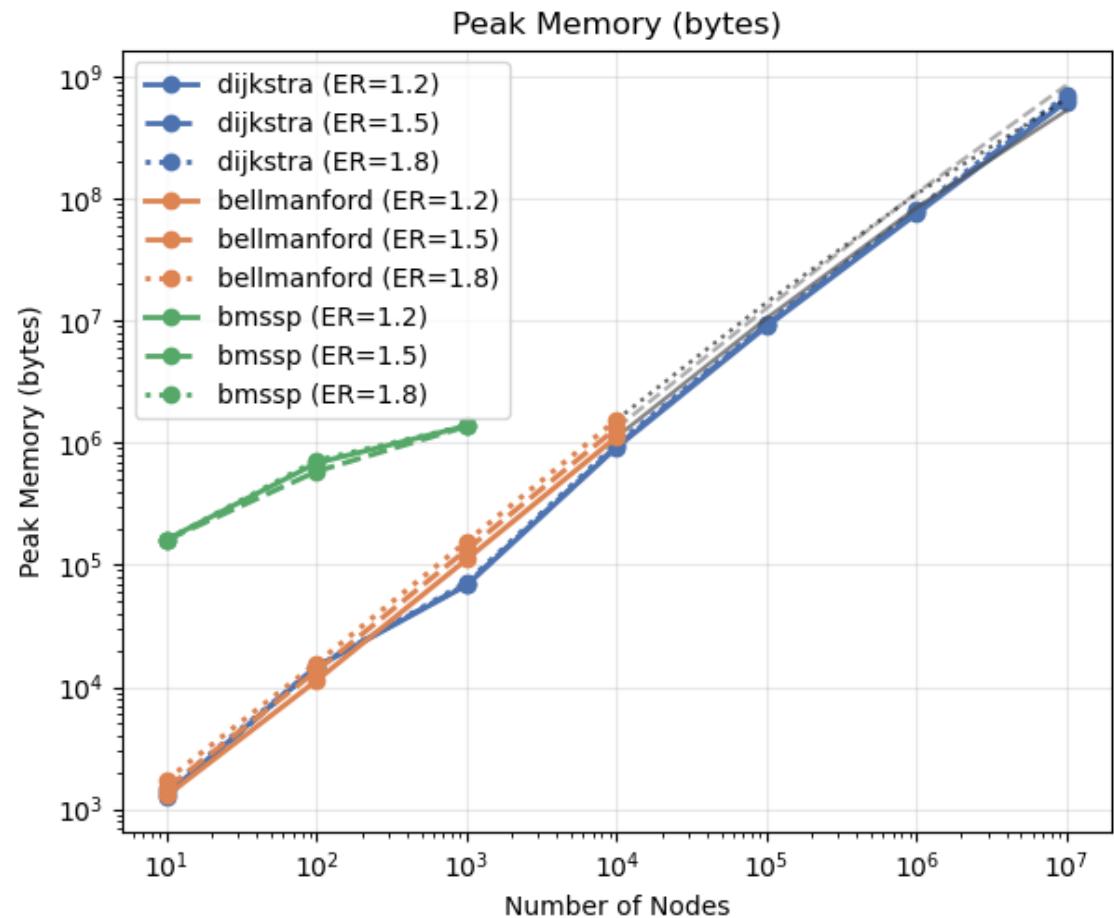
Results: Time Complexity

- As expected, Dijkstra's algorithm was significantly faster than the other two algorithms.
- BMSSP took considerably longer but still exhibited a quasi-polynomial growth pattern.
- Although BMSSP is theoretically superior to Dijkstra's, its advantages may be hard to justify for graph sizes under billions of nodes.



Results: Space Complexity

- Dijkstra and Bellman-Ford showed similar peak memory usage across graphs.
- BMSSP's memory footprint was much larger due to repeated data structure reinitialization and recursion.



DEMO TIME